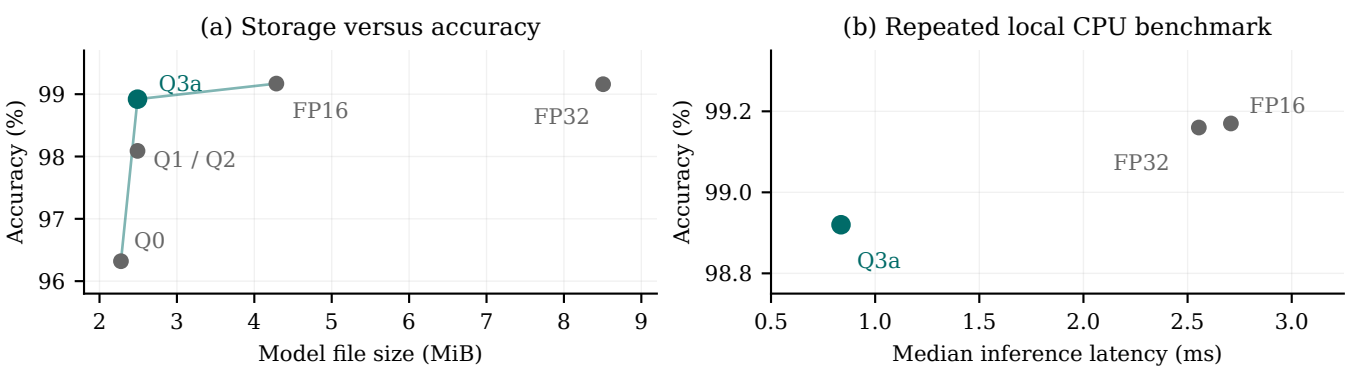
(a) Storage versus accuracy
Accuracy (%)
99
98
97
96
Q3a
FP16
FP32
Q1 / Q2
Q0
2
3
4
5
6
7
8
9
Model file size (MiB)
(b) Repeated local CPU benchmark
Accuracy (%)
99.2
99.0
98.8
FP16
FP32
Q3a
0.5
1.0
1.5
2.0
2.5
3.0
Median inference latency (ms)